



Fig. 2. Ablation study: ASR, RD, CI and AR across defense configurations (dataset: 10 queries $\times$ 6 attack strategies; generator: GPT-3.5-turbo)


Attack Success Rate (ASR)




Robustness Degradation (RD)



Context Integrity (CI)



Answer Relevance (AR)



Single-layer / No defense
 All defenses (proposed)
 No-defense baseline